

# Topic 3 Lecture - Introduction to Computer Vision and Machine Learning



# Aims of the Session

- Learn the basics on how images are imported to our computers and converted into data



# Resources for the Lecture

- Introduction to Computing and Programming in Python: A Multimedia Approach.  
Mark Guzdial, Barbara Ericson. Pearson, 2016.
- Various others mentioned throughout the lecture!

How does an image look *digitally*?

# How does an image look *digitally*?

- These are the main **compression algorithms** used to store images:

# How does an image look *digitally*?

- These are the main **compression algorithms** used to store images:





# Compression Algorithms

# Compression Algorithms

- The *art* of compression algorithms is in **quantisation**!

# Compression Algorithms

- The *art* of compression algorithms is in **quantisation**!
- Best algorithms are the ones that achieve best visual quality with reduced size.

# Compression Algorithms

- The *art* of compression algorithms is in **quantisation**!
- Best algorithms are the ones that achieve best visual quality with reduced size.



8.9M



68.34K

- The importance of compression

- The importance of compression

```
In [1]: import warnings
warnings.filterwarnings('ignore')
from IPython.display import HTML
HTML('<iframe width="560" height="315" src="https://www.youtube.com/embed/NMkz
```

Out[1]:

The first movie with CGI



- Fortunately, this is **NOT** our problem in this module!

- Fortunately, this is **NOT** our problem in this module!
- We are going to work with images in *simpler* ways





# Images as arrays/matrices

- Using the `numpy` module

# Images as arrays/matrices

- Using the `numpy` module
- Complementing by using the `OpenCV` module, which will let us import and manipulate images

- When we import an image, the first thing we will get is a **bitmap**

- When we import an image, the first thing we will get is a **bitmap**



- When we import an image, the first thing we will get is a **bitmap**



- Each pixel will be represented as a value within an  $n \times m$  matrix

Grayscale Images

# Grayscale Images

- A 2D grid of pixels



## Grayscale Images

- A 2D grid of pixels
- Two ways to represent them:

1. Standard: from 0 (black) to 255 (white) with 254 gray values in between.

1. Standard: from 0 (black) to 255 (white) with 254 gray values in between.




187	183	174	168	160	152	128	181	172	161	155	166				
195	182	163	74	75	62	33	17	110	230	180	154				
180	180	80	14	34	6	10	33	48	106	199	181				
206	106	5	124	131	111	120	204	166	15	56	180				
194	68	137	251	237	238	236	228	227	67	71	201				
172	106	207	233	239	214	220	239	228	94	74	206				
188	88	179	208	189	216	211	198	136	75	20	169				
189	97	166	84	10	168	134	11	31	62	22	148				
199	168	181	192	198	227	178	143	182	106	36	190				
205	174	155	252	236	231	148	179	238	49	95	234				
190	216	116	148	236	187	86	158	79	38	218	241				
190	224	147	108	227	210	127	103	36	101	255	224				
190	214	173	64	103	143	94	90	2	109	249	215				
187	196	226	75	1	81	47	0	6	217	255	211				
183	202	237	145	0	0	12	108	200	136	243	236				
195	206	123	207	177	121	123	200	179	13	96	218				

187	183	174	168	160	152	128	181	172	161	155	166				
195	182	163	74	75	62	33	17	110	230	180	154				
180	180	80	14	34	6	10	33	48	106	199	181				
206	106	5	124	131	111	120	204	166	15	56	180				
194	68	137	251	237	238	236	228	227	67	71	201				
172	106	207	233	239	214	220	239	228	94	74	206				
188	88	179	208	189	216	211	198	136	75	20	169				
189	97	166	84	10	168	134	11	31	62	22	148				
199	168	181	192	198	227	178	143	182	106	36	190				
205	174	155	252	236	231	148	179	238	49	95	234				
190	216	116	148	236	187	86	158	79	38	218	241				
190	224	147	108	227	210	127	103	36	101	255	224				
190	214	173	64	103	143	94	90	2	109	249	215				
187	196	226	75	1	81	47	0	6	217	255	211				
183	202	237	145	0	0	12	108	200	136	243	236				
195	206	123	207	177	121	123	200	179	13	96	218				

1. Normalised: from 0 (black) to 1 (white) with "infinite" gray values in between.


1. Normalised: from 0 (black) to 1 (white) with "infinite" gray values in between.



**IS IT POSSIBLE TO CONVERT BETWEEN STANDARD  $\leftrightarrow$  NORMALISED?**

# Colour Images

# Colour Images





- Each pixel has three channels : *red*, *green* and *blue*

- Each pixel has three channels : *red*, *green* and *blue*
- Images with colour are often called *RGB* images

Option 1

### Option 1

- If a colour image is imported, a matrix will be produced, this time with three values per pixel instead of one

### Option 1

- If a colour image is imported, a matrix will be produced, this time with three values per pixel instead of one
- The three values will be stored in a tuple

	0	1	2	3	4	5	6	7	8	9
0										
1										
2										
3										
4										
5										
6										

RGB (218, 150, 149)

R = 11011010

G = 10010110

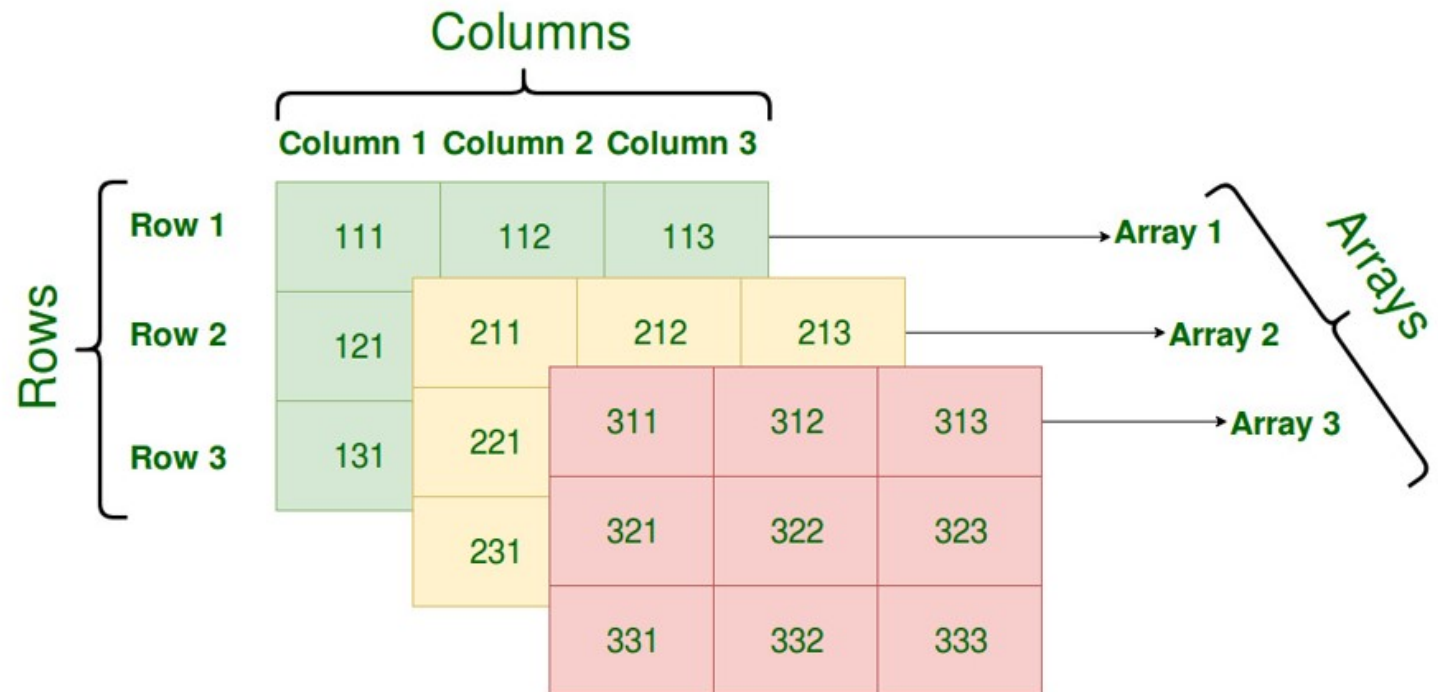
B = 10010101

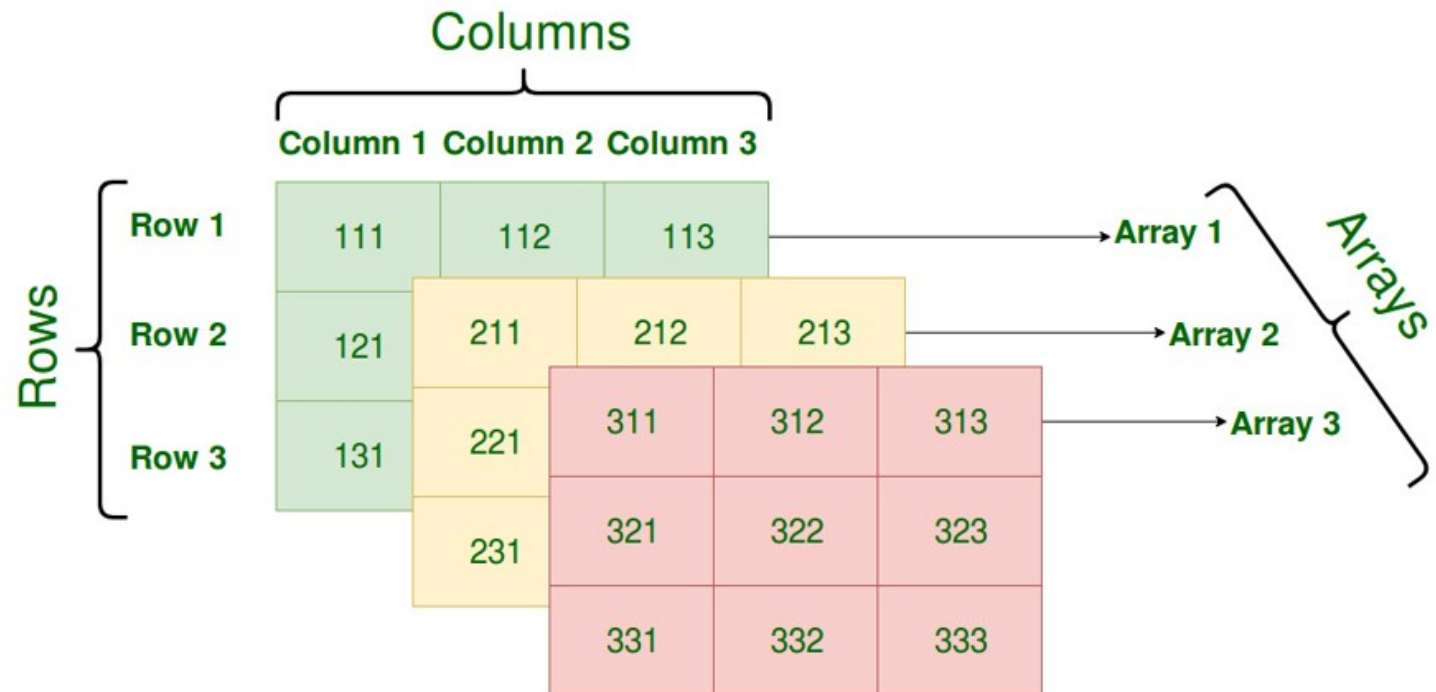
Option 2

## Option 2

- When importing a colour image in OpenCV , a 3D array will be produced, with the third dimension representing the three channels







- Advantage of option 2: Faster to do calculations and transformations

**HOW MANY COLOURS CAN BE REPRESENTED USING THIS STANDARD?**

**HOW MANY COLOURS CAN BE REPRESENTED USING THIS STANDARD?**

**CAN RGB BE NORMALISED?**

**HOW MANY COLOURS CAN BE REPRESENTED USING THIS STANDARD?**

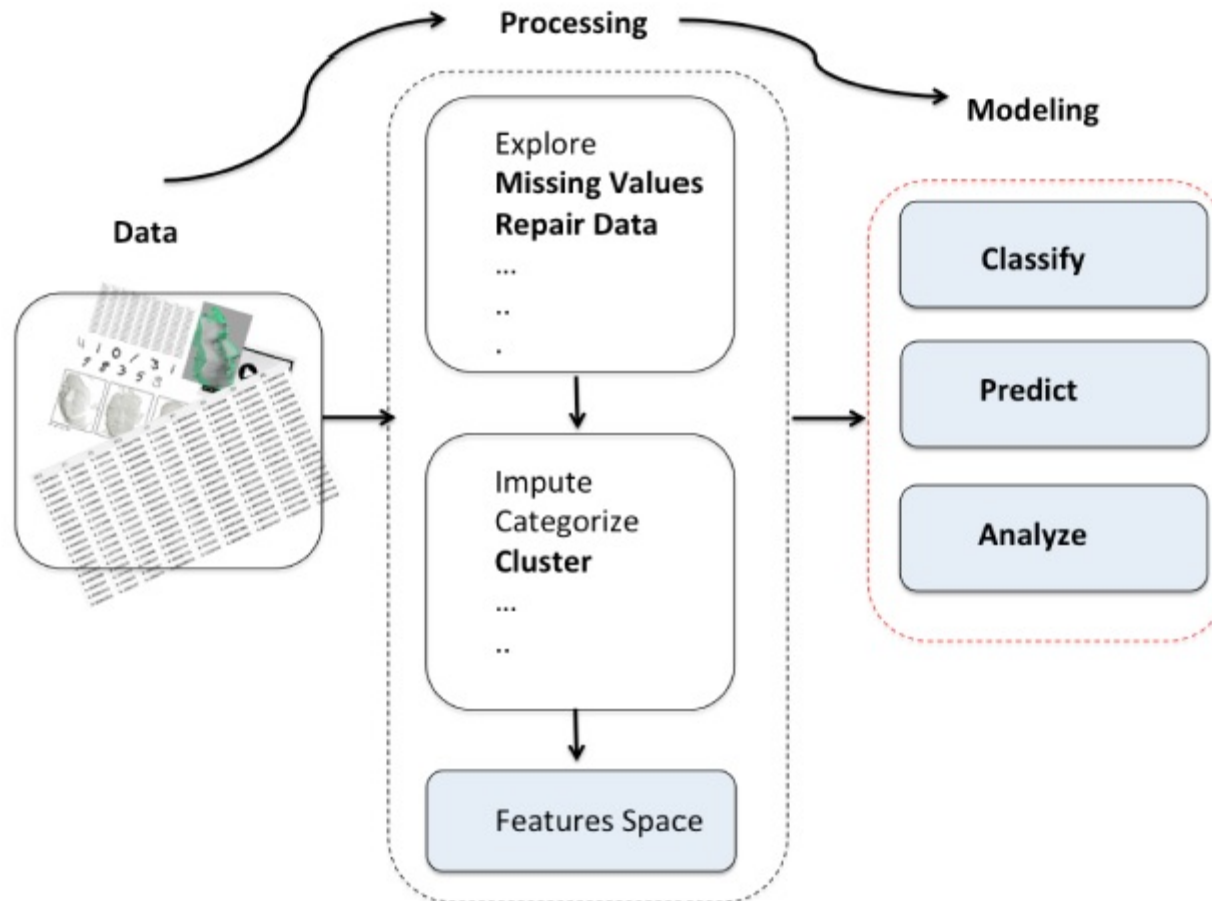
**CAN RGB BE NORMALISED?**

**ARE THERE ANY OTHER STANDARDS THAT CAN REPRESENT MORE COLOURS?**



# Machine Learning

- Machine learning: An automatic function that maps  $x \rightarrow y$  based on the input data





# Types of Machine Learning

# Types of Machine Learning

Supervised Learning

# Types of Machine Learning

## Supervised Learning

- Aims to learn a function that, given a sample of data and desired outputs, approximates a function that maps inputs to output

# Types of Machine Learning

## Supervised Learning

- Aims to learn a function that, given a sample of data and desired outputs, approximates a function that maps inputs to output
- Done in the context of **classification** (when mapping input to output label) or **regression** (when mapping input to continuous output)

# Types of Machine Learning

## Supervised Learning

- Aims to learn a function that, given a sample of data and desired outputs, approximates a function that maps inputs to output
- Done in the context of **classification** (when mapping input to output label) or **regression** (when mapping input to continuous output)
- The "correct" output will be deduced from the training data, therefore the model requires a reliable base

## Examples of Supervised Learning Algorithms

## Regression


## Regression

- The simplest ML out there! A line (or curve) that adapts to the data points and tries to do a prediction based on the existing data and how it adapts to a mathematical function



## Regression


- The simplest ML out there! A line (or curve) that adapts to the data points and tries to do a prediction based on the existing data and how it adapts to a mathematical function



Source

- There's a version called logistic regression which uses a probability curve, which in turn can be adapted for binary classification!

- There's a version called logistic regression which uses a probability curve, which in turn can be adapted for binary classification!



Source

- Also very simple code to implement



- Also very simple code to implement

```
In [2]: # Source: https://www.w3schools.com/python/python\_ml\_linear\_regression.asp

import matplotlib.pyplot as plt
from scipy import stats


x = [5,7,8,7,2,17,2,9,4,11,12,9,6]
y = [99,86,87,88,111,86,103,87,94,78,77,85,86]

# This is the lin reg, one line of code!
slope, intercept, r, p, std_err = stats.linregress(x, y)

def myfunc(x):
    return slope * x + intercept

mymodel = list(map(myfunc, x))

plt.scatter(x, y, color='red')
plt.plot(x, mymodel)
plt.show()
```



K Nearest Neighbours



K Nearest Neighbours


- Classify by a majority vote of neighbours

### K Nearest Neighbours

- Classify by a majority vote of neighbours
- Advantages: Simple to implement, robust to noisy training data, and effective if training data is large

### K Nearest Neighbours

- Classify by a majority vote of neighbours
- Advantages: Simple to implement, robust to noisy training data, and effective if training data is large
- Disadvantages: Need to determine the value of K computation cost is high



In [3]: *# Source: [https://www.w3schools.com/python/python\\_ml\\_knn.asp](https://www.w3schools.com/python/python_ml_knn.asp)*


```
import matplotlib.pyplot as plt
from sklearn.neighbors import KNeighborsClassifier

x = [4, 5, 10, 4, 3, 11, 14, 8, 10, 12]
y = [21, 19, 24, 17, 16, 25, 24, 22, 21, 21]
classes = [0, 0, 1, 0, 0, 1, 1, 0, 1, 1]
data = list(zip(x, y))

# Two lines of code to declare and "fit" the model!
knn = KNeighborsClassifier(n_neighbors=1)
knn.fit(data, classes)

# Test it
new_x = 8
new_y = 21

new_point = [(new_x, new_y)]
prediction = knn.predict(new_point)
plt.scatter(x + [new_x], y + [new_y], c=classes + [prediction[0]])
plt.text(x=new_x-1.7, y=new_y-0.7, s=f"new point, class: {prediction[0]}")
plt.show()
```



Decision Tree/Random Forest

Decision Tree/Random Forest

- Segment the predictor space into multiple regions



### Decision Tree/Random Forest

- Segment the predictor space into multiple regions
- Each region has only a subset of the training dataset

### Decision Tree/Random Forest

- Segment the predictor space into multiple regions
- Each region has only a subset of the training dataset
- High variance → Small changes in the training data can give an entirely different decision tree model

Is a Person Fit?

Age < 30 ?

Yes?

No?

Eat's a lot  
of pizzas?

Exercises in  
the morning?

Yes?

No?

Yes?

No?

Unfit!

Fit

Fit

Unfit!

In [4]: *# Source: <https://scikit-learn.org/stable/modules/tree.html>*

```
from sklearn.datasets import load_iris
from sklearn import tree

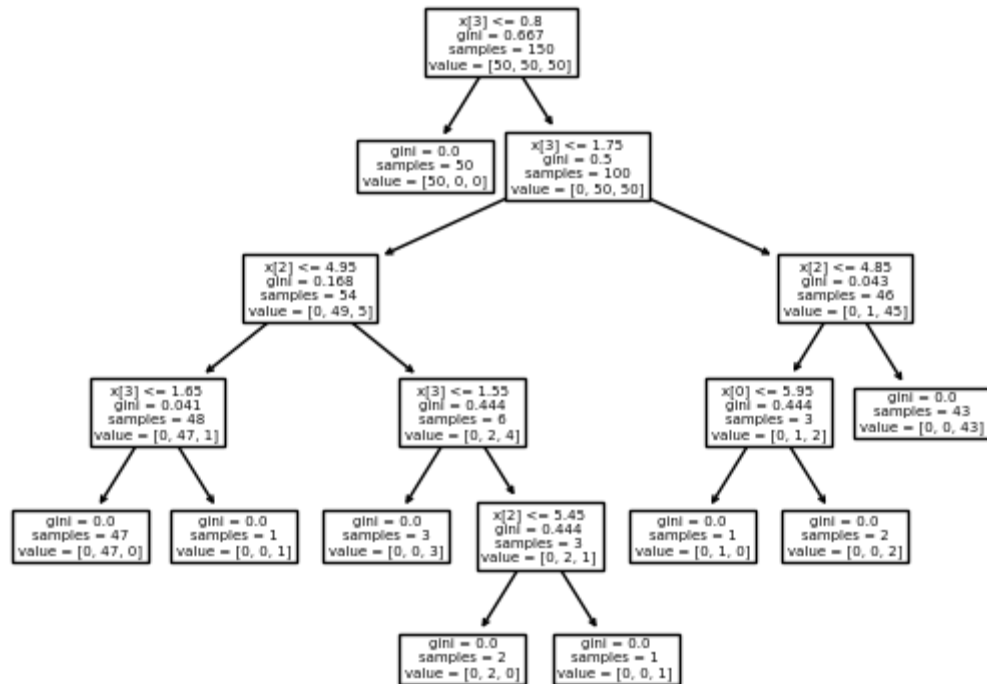
iris = load_iris()
X, y = iris.data, iris.target

## Again, classifier is just two lines of code
clf = tree.DecisionTreeClassifier()
clf = clf.fit(X, y)

# You can even visualise the tree
tree.plot_tree(clf)
```

Out[4]: [Text(0.5, 0.9166666666666666, 'x[3] <= 0.8\ngini = 0.667\nsamples = 150\nvalue = [50, 50, 50]'),  
Text(0.4230769230769231, 0.75, 'gini = 0.0\nsamples = 50\nvalue = [50, 0, 0]'),  
Text(0.5769230769230769, 0.75, 'x[3] <= 1.75\ngini = 0.5\nsamples = 100\nvalue = [0, 50, 50]'),  
Text(0.3076923076923077, 0.5833333333333333, 'x[2] <= 4.95\ngini = 0.168\nsamples = 54\nvalue = [0, 49, 5]'),  
Text(0.15384615384615385, 0.4166666666666667, 'x[3] <= 1.65\ngini = 0.041\nsamples = 48\nvalue = [0, 47, 1]'),  
Text(0.07692307692307693, 0.25, 'gini = 0.0\nsamples = 47\nvalue = [0, 47, 0]'),  
Text(0.23076923076923078, 0.25, 'gini = 0.0\nsamples = 1\nvalue = [0, 0, 1]'),  
Text(0.46153846153846156, 0.4166666666666667, 'x[3] <= 1.55\ngini =

```
0.444\nsamples = 6\nvalue = [0, 2, 4]'),  
  Text(0.38461538461538464, 0.25, 'gini = 0.0\nsamples = 3\nvalue =  
[0, 0, 3]'),  
  Text(0.5384615384615384, 0.25, 'x[2] <= 5.45\ngini = 0.444\nsamples  
= 3\nvalue = [0, 2, 1]'),  
  Text(0.46153846153846156, 0.08333333333333333, 'gini = 0.0\nsamples  
= 2\nvalue = [0, 2, 0]'),  
  Text(0.6153846153846154, 0.08333333333333333, 'gini = 0.0\nsamples =  
1\nvalue = [0, 0, 1]'),  
  Text(0.8461538461538461, 0.5833333333333334, 'x[2] <= 4.85\ngini =  
0.043\nsamples = 46\nvalue = [0, 1, 45]'),  
  Text(0.7692307692307693, 0.4166666666666667, 'x[0] <= 5.95\ngini =  
0.444\nsamples = 3\nvalue = [0, 1, 2]'),  
  Text(0.6923076923076923, 0.25, 'gini = 0.0\nsamples = 1\nvalue = [0,  
1, 0]'),  
  Text(0.8461538461538461, 0.25, 'gini = 0.0\nsamples = 2\nvalue = [0,  
0, 2]'),  
  Text(0.9230769230769231, 0.4166666666666667, 'gini = 0.0\nsamples =  
43\nvalue = [0, 0, 43]')]
```



Naive Bayes

## Naive Bayes

- Probabilistic classifier inspired by the Bayes theorem, assumes attributes are conditionally independent



## Naive Bayes

- Probabilistic classifier inspired by the Bayes theorem, assumes attributes are conditionally independent
- Advantages: small amount of training data required, extremely fast

## Naive Bayes

- Probabilistic classifier inspired by the Bayes theorem, assumes attributes are conditionally independent
- Advantages: small amount of training data required, extremely fast
- Disadvantages: zero probability problem, if the conditional probability is zero for a particular attribute...



Posterior

Attributes likelihoods

$$P(y|x_1, \dots, x_n) = \frac{P(x_1|y)P(x_2|y)\dots P(x_n|y)P(y)}{P(x_1)P(x_2)\dots P(x_n)}$$

Predictor prior

Class prior



In [5]: *# Credit: <https://www.kaggle.com/code/nizamudma/iris-data-classification-using>*

```
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB

#split data for train and test
x_train,x_test,y_train,y_test=train_test_split(X,y,test_size=0.15)

# No surprise, two lines
bc=GaussianNB()
bc.fit(x_train,y_train)

# Predict samples from the test set
print(bc.predict(x_test))
```

```
[0 1 2 1 1 0 0 2 0 2 0 0 0 2 1 1 1 0 0 2 2 0 0]
```

Support Vector Machine (SVM)

## Support Vector Machine (SVM)

- Discriminative classifier defined by a separating hyperplane



## Support Vector Machine (SVM)

- Discriminative classifier defined by a separating hyperplane
- Tuning parameters in SVM classifier
  - Kernel - transformation method, e.g. Polynomial and exponential kernels
  - Regularisation - how much to avoid misclassifying each training example
  - Gamma - how far the influence of a single training example reaches, high gamma  $\rightarrow$  only nearby examples






## Support Vector Machine (SVM)

- Discriminative classifier defined by a separating hyperplane
- Tuning parameters in SVM classifier
  - Kernel - transformation method, e.g. Polynomial and exponential kernels
  - Regularisation - how much to avoid misclassifying each training example
  - Gamma - how far the influence of a single training example reaches, high gamma  $\rightarrow$  only nearby examples
- A margin in SVM is a separation of line to the closest class points
  - A good margin is one where this separation is larger for both the classes



Transform  
 $Z = x^2 + y^2$







- We will do this one in Lab Topic 5! But trust me, again it is two lines of code...

Neural Network (NN)

Neural Network (NN)

We will talk **EXTENSIVELY** about them later on



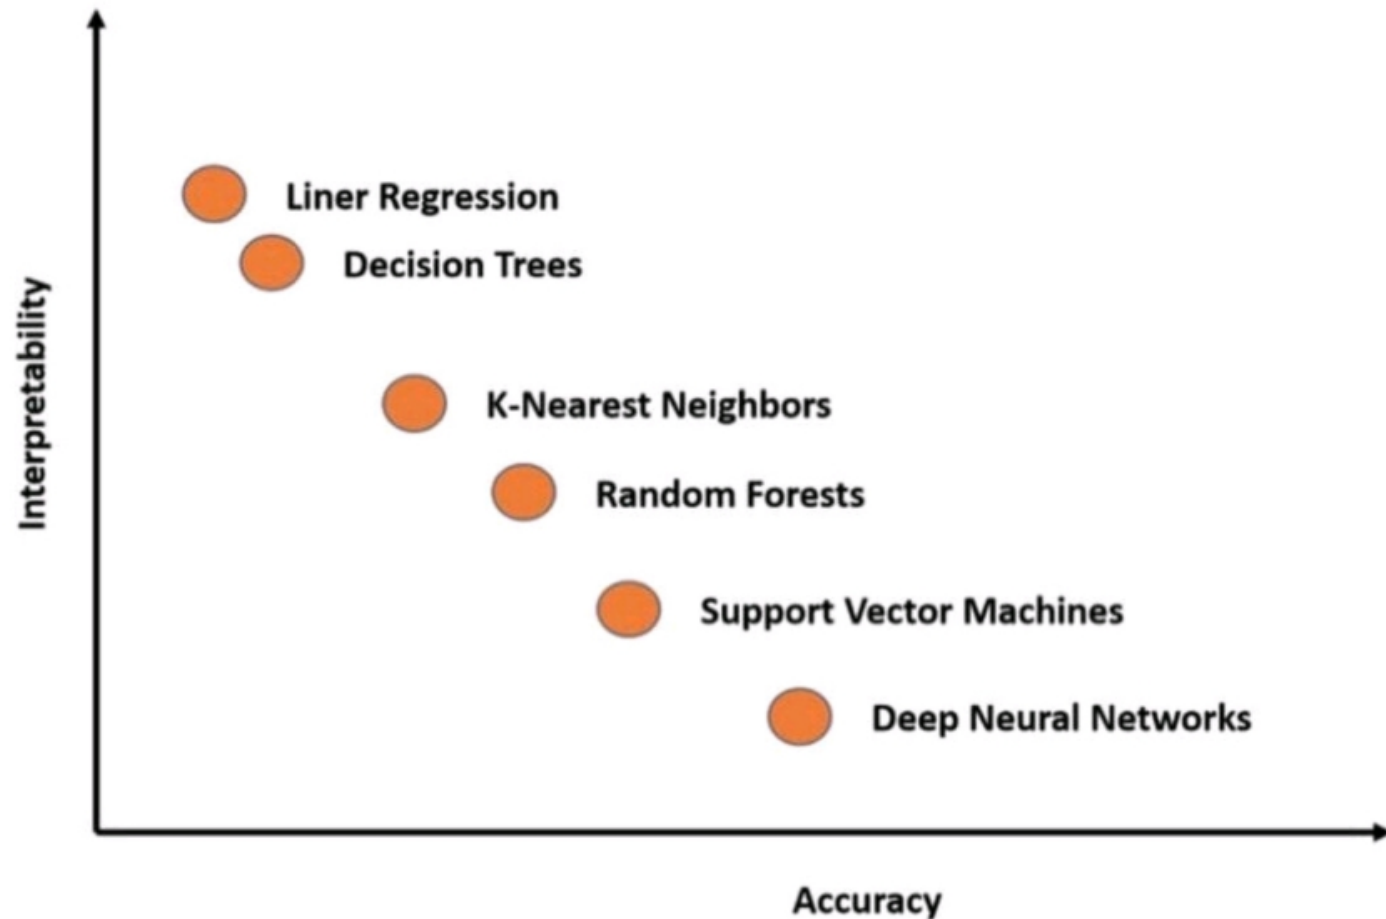




**Accuracy**

**and**

**Interpretability**





## Semi-supervised Learning

## Semi-supervised Learning

- Aims to label unlabelled data points using knowledge learned from a small number of labelled data points

## Semi-supervised Learning

- Aims to label unlabelled data points using knowledge learned from a small number of labelled data points
- Learning with both unlabelled and labeled data points

## Semi-supervised Learning

- Aims to label unlabelled data points using knowledge learned from a small number of labelled data points
- Learning with both unlabelled and labeled data points
- Used when large amounts of data are costly to label

## Assumptions of Semi-supervised Learning

## Assumptions of Semi-supervised Learning

- **Continuity** : Data points that are "close" have a common label



### Assumptions of Semi-supervised Learning

- **Continuity** : Data points that are "close" have a common label
- **Cluster** : Data naturally forms discrete clusters, most common to share label

### Assumptions of Semi-supervised Learning

- **Continuity** : Data points that are "close" have a common label
- **Cluster** : Data naturally forms discrete clusters, most common to share label
- **Manifold** : Data lies in a lower dimensional space than the input space

## Examples of Semi-supervised Learning Algorithms

## Examples of Semi-supervised Learning Algorithms

Transductive SVM

## Examples of Semi-supervised Learning Algorithms

Transductive SVM

Label Propagation

## Unsupervised Learning

## Unsupervised Learning

- Does not have (or need) any labelled outputs, so its goal is to infer the natural structure present within a set of data points

## Unsupervised Learning

- Does not have (or need) any labelled outputs, so its goal is to infer the natural structure present within a set of data points
- Finds inherent patterns of data



## Unsupervised Learning

- Does not have (or need) any labelled outputs, so its goal is to infer the natural structure present within a set of data points
- Finds inherent patterns of data
- Most common tasks: **clustering** and **exploratory data analysis**

## Examples of Unsupervised Learning Algorithms

## Examples of Unsupervised Learning Algorithms

K-means Clustering

## Examples of Unsupervised Learning Algorithms

K-means Clustering

Principal Component Analysis (PCA)

## Examples of Unsupervised Learning Algorithms

K-means Clustering

Principal Component Analysis (PCA)

Autoencoders



## A simpler explanation of what I just said

```
In [6]: import warnings
warnings.filterwarnings('ignore')
from IPython.display import HTML
HTML('<iframe width="560" height="315" src="https://www.youtube.com/embed/R90I
```

Out[6]:

How AIs, like ChatGPT, Learn



As you can see, we can do more things beyond classification, such as **detecting** where the object lies in the image, **recognise** it from a pool of similar objects, **segment** where it is or even **track** it in a video feed!



